



UNIVERSITY OF SAN FRANCISCO  
CHANGE THE WORLD FROM HERE

---

# Concurrency

Sami Rollins



# Concurrency

---

- What is concurrency?

# Concurrency

---

- **What is concurrency?**

- Executing multiple tasks simultaneously
  - Processing multiple HTTP requests at the same time

- **Why concurrency?**

- Improve performance
  - Take advantage of multiple cores
  - Do useful computation when blocked on i/o
- Allow background activities

## GrepLite

---

searchterm: computing

directory: /Users/srollins/GrepTool/myfiles



file 1: <count>

file 2: <count>

...

file n: <count>

## GrepLite Algorithm - Sequential

---

time



create new FileProcessor  
processFile()

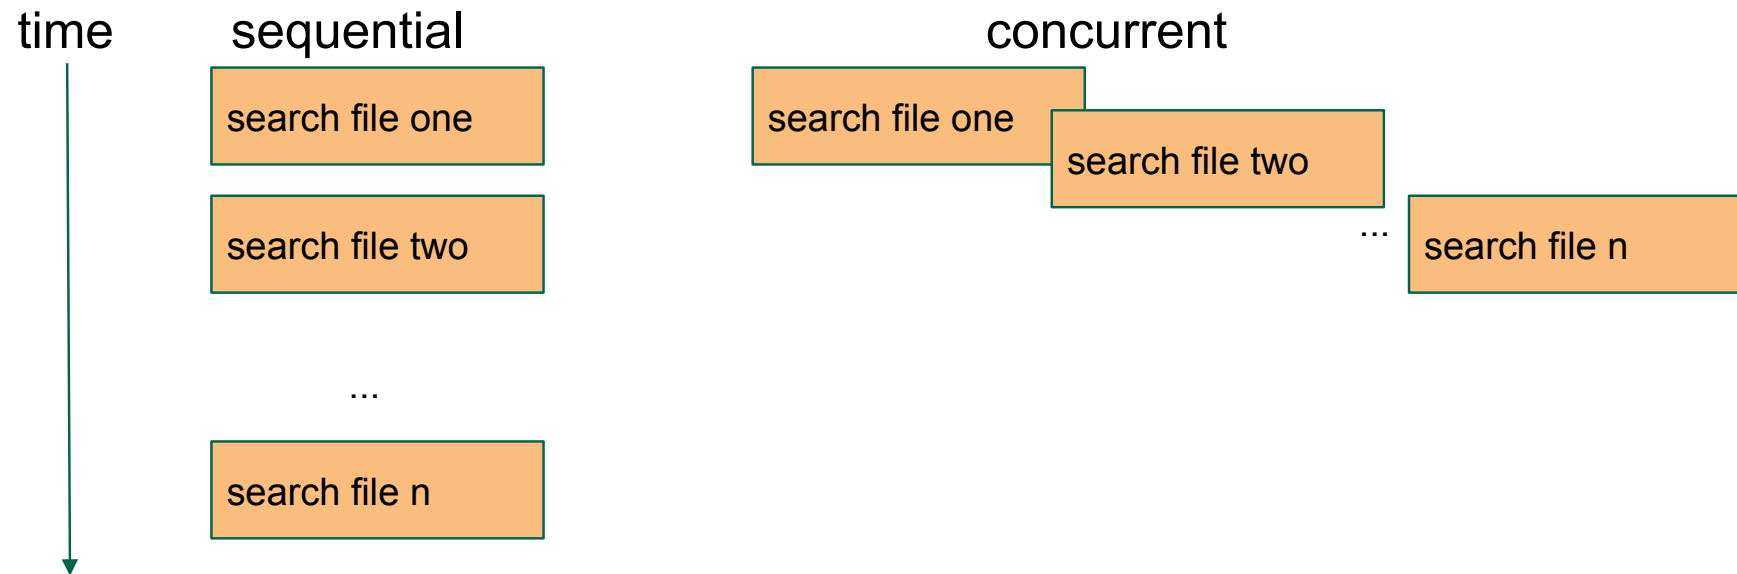
for each line in file  
    if line contains term  
        linecount++  
print result

create new FileProcessor  
processFile()

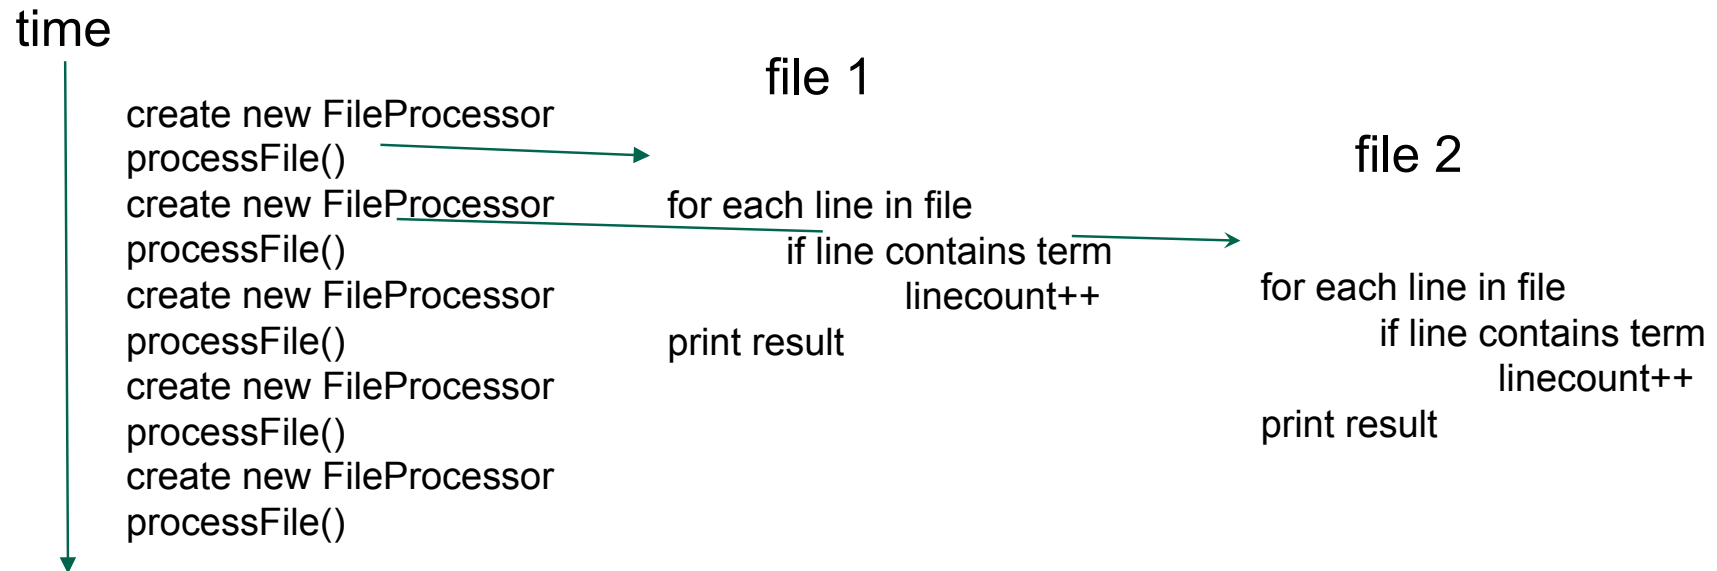
for each line in file  
    if line contains term  
        linecount++  
print result

## Sequential versus Concurrent

---




## GrepLite Algorithm - Concurrent



## Threads

---

- Programming abstraction that allow *lightweight processes* within an application





## Challenges

---

- **Order of execution**
  - Race conditions
- **Inter-thread communication**
- **Shared data and thread safety**
- **Overhead**

## Threads and sharing

---

### Atomic operations

Suppose two threads operate on the same data  
a initially has the value of 4

Thread 1:

$a = a + 1;$

Thread 2:

$a = a * 2;$

What are the possible resulting values for a?

10 ( $4 + 1 = 5$ ,  $5 * 2 = 10$ )

9 ( $4 * 2 = 8$ ,  $8 + 1 = 9$ )

how about 5 or 8?

## Threads and sharing

---

$a = a + 1$

1. read value of  $a$
2. set  $a$  to  $a + 1$

$a = a * 2$

3. read value of  $a$
4. set  $a$  to  $a * 2$

Option 1: 1, 2, 3, 4 --  $a$  will be 10

Option 2: 3, 4, 1, 2 --  $a$  will be 9

Option 3: 1, 3, 2, 4 --  $a$  will be 8

Option 4: 1, 3, 4, 2 --  $a$  will be 5

BUT, options 3 and 4 should never occur

## Atomic operations

---

*atomic* operations execute without interference from other operations

request a *lock* before an operation, release it when the operation is complete

`synchronized` is a Java keyword that acquires a lock on an object

`synchronized` may be used for methods or blocks of statements

synchronized methods may not interleave

## Thread pool

---

There is an overhead associated with creating and destroying threads  
Each operation requires a new thread

A solution is to create a *pool* of a fixed number of threads. As tasks arrive, a thread from the pool is used to execute the task.

array of threads



queue of tasks



## Communication

---

Threads *wait* for new work

When new work is inserted, the queue must *notify* the threads that new work is available

`wait` and `notify` are Java keywords.

`wait` can be called on any object. The thread that calls `wait` suspends execution until it is *woken up* with a `notify` message called on the same object.

In most cases, `notifyAll` is used instead of `notify`. `notifyAll` will wake up every thread that has called `wait` on the given object.