



UNIVERSITY OF SAN FRANCISCO
CHANGE THE WORLD FROM HERE

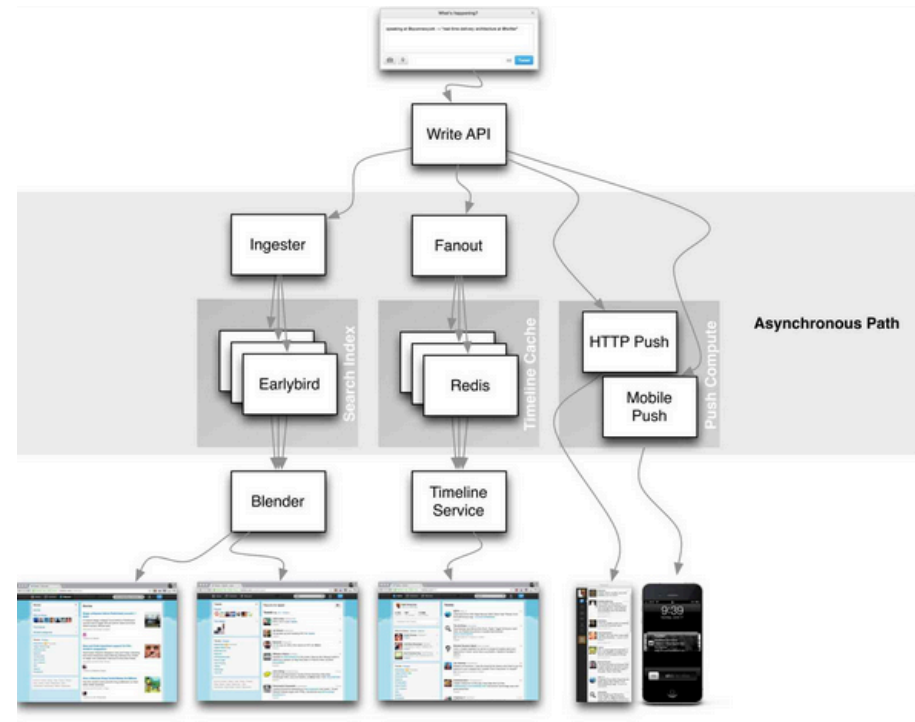
DSD Overview

Sami Rollins




Why study distributed systems?

- You use them everyday!
- Scale is the name of the game
- **Distributed systems are complex**
 - Understanding what is under the hood is important

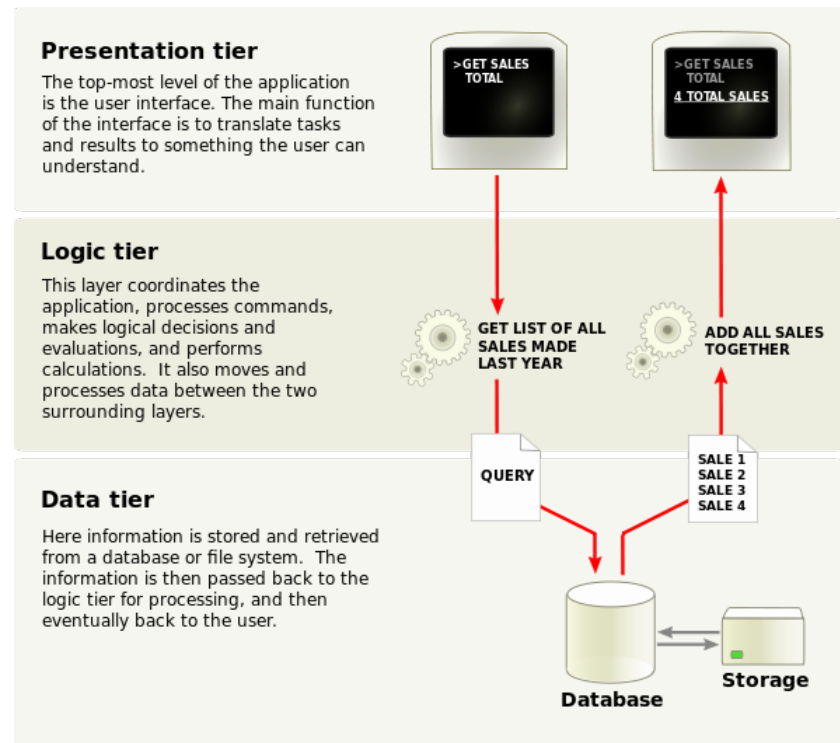


History lesson



- **Client/server or n-tier Architecture**

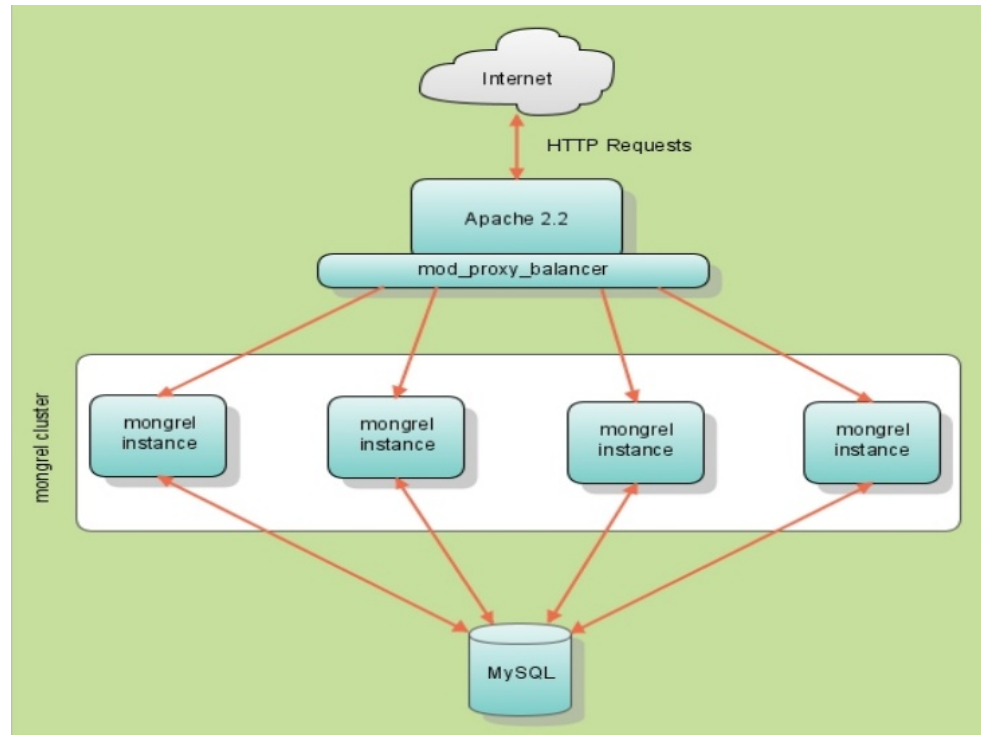
- Limited scalability
- Database often a bottleneck



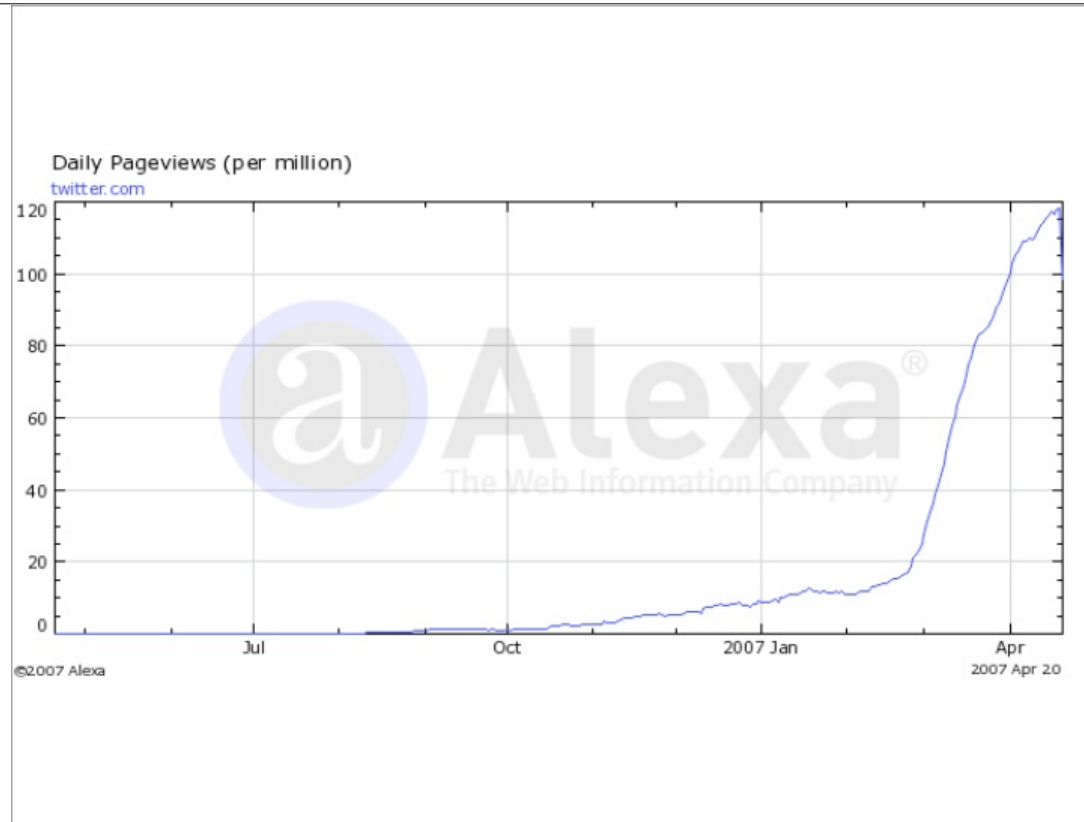
Twitter case study

- **Original Twitter architecture**

- Ruby on Rails
- MySQL
 - 1 server
 - 2400 requests/second
- Mongrel (web server)
- Memcached

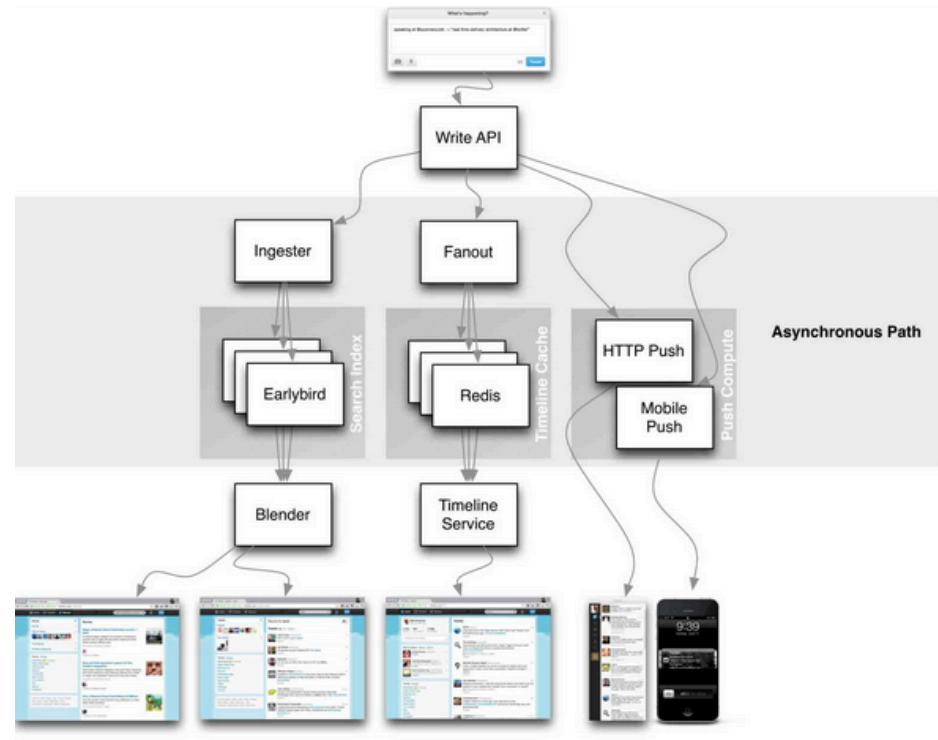


Twitter case study



Twitter case study

- **Optimize reads**
 - insert
 - for each follower, insert new tweet in redis cache
- **Service-oriented architecture**
 - System is a set of loosely coupled services





Non-web distributed systems?

- **Twitter, Pinterest, Google, Etsy**
 - All are web-based applications
- **What about other kinds of applications?**
- **Characteristics**
 - Concurrency
 - Independent failures
 - No global clock
 - Heterogeneity
 - High latency communication
 - System composed of computers spread over a large geographic area, maybe under different administrative domains


Other distributed systems

- **Peer-to-peer**
 - File sharing/distribution
 - Gnutella?
 - Bit torrent
 - Others?
- **MMOG**
- **DNS**
- **CDNs**
 - Akamai
 - Limelight
 - Netflix Open Connect

Architectures



Architectures



Bit torrent - <http://en.wikipedia.org/wiki/Peer-to-peer>

Main challenges?

- What do you think are the biggest challenges in building a distributed system?

Challenges

- **Heterogeneity**
 - different networks, OSs, architectures
 - communication protocols must be carefully defined
- **Openness**
 - Can new services/components be easily added?
- **Security**
 - confidentiality, protection against corruption, resistant to attack
- **Scalability**
 - tolerates increase in users and/or resources

Challenges

- **Failure handling**
 - Can the system detect, mask, tolerate, recover from failures?
- **Concurrency**
 - able to handle multiple requests simultaneously
- **Transparency**
 - local and remote resources accessible in the same way

Principles

- **Availability**
 - Enough resources to handle all traffic
 - Failure is handled gracefully
- **Performance**
 - Fast response time
 - Sufficient hardware
 - Geographic placement


Principles

- **Reliability**
 - Fault tolerance
 - Consistency
 - Eventual consistency
 - Strong consistency
- **Scalability**
 - Increasing resources, supporting more users
- **Manageability**
 - System updates, failures
- **Cost**

CAP theorem – Pick two

- **Consistency**
 - All nodes see the same data at the same time
- **Availability**
 - A guarantee that every request receives a response about whether it was successful or failed
- **Partition tolerance**
 - The system continues to operate despite arbitrary message loss or failure of part of the system

Pieces of the puzzle



Design decisions

- **Load balancing algorithms**
 - DNS round robin
 - Random
 - CPU/Memory utilization
- **Service-oriented Architecture (SOA)**
 - Replication versus partitioning of functionality
- **Caching**
 - Policies for populating cache and evicting data
 - Memcached
 - Redis

Design decisions

- **Data**
 - Replicated versus partitioned
- Asynchronous processing
 - Log requests for “big data” processing, transformation, etc
 - Kafka